

# Decision Tree Model

# Decision Tree

- In principle, there are **exponentially many decision trees** that can be constructed from a given set of attributes.
- Finding the optimal tree (in terms of performance) is computationally **infeasible** because of the exponential size of the search space.
- Efficient algorithms have been developed to induce a reasonably **accurate** although **suboptimal** decision tree in a reasonable amount of time.
- These algorithms usually employ a **greedy strategy** that grows a decision tree by making a series of locally optimum decision about which attribute to use for partitioning the data.
- One such algorithm is Hunt's algorithm, which is the basis of many decision tree induction algorithms, including ID3, C4.5, and CART.

# Hunt's Algorithm

- In this algorithm, a decision tree is grown in a recursive fashion by partitioning the training records into successively purer subsets.
- Let  $D_t$  be the set of training records that are associated with node  $t$  and  $y = \{y_1, y_2, \dots, y_c\}$  be the class labels. The following is a recursive definition of Hunt's algorithm:
  - Step 1: If all the records in  $D_t$  belong to the same class  $y_t$ , then  $t$  is a leaf node labelled as  $y_t$ .
  - Step 2: If  $D_t$  contains records that belong to more than one class, an **attribute test condition** is selected to partition the records into smaller subsets. A child node is created for each outcome of the test condition and the records in  $D_t$  are distributed to the children based on the outcomes.
  - The algorithm is then recursively applied to each child node.

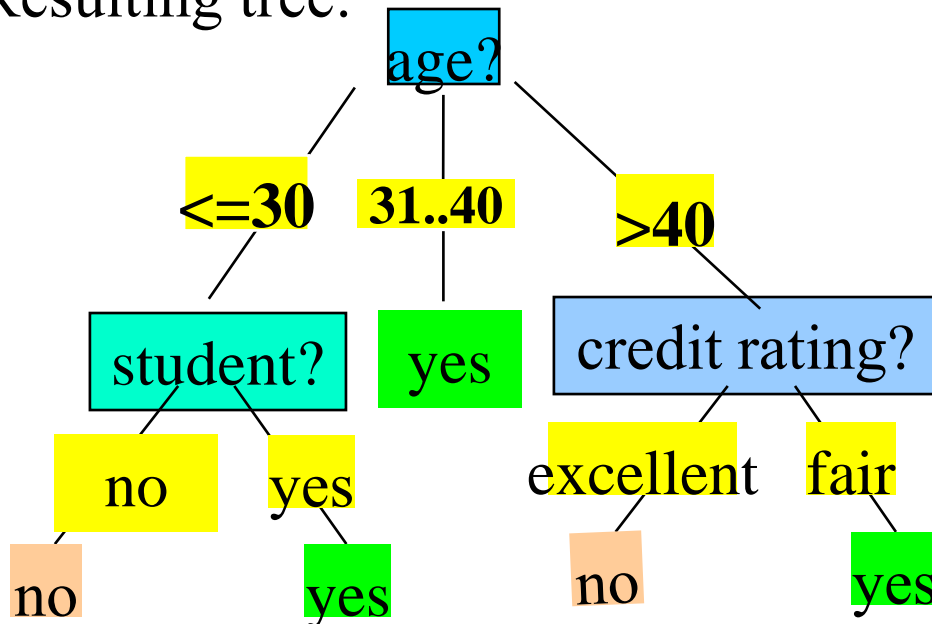
# Decision Tree Induction: ID3

- ❑ Training data set: Buys\_computer
- ❑ Decision tree model is constructed using the **Quinlan ID3** algorithm.
- ❑ ID3 refers to Iterative Dichotomizer 3, and is developed by Ross Quinlan
- ❑ It iteratively dichotomizes (divides) the data using a top down greedy approach.
- ❑ The algorithm optimizes locally at each iteration.

age	income	student	credit_rating	buys_computer
<=30	high	no	fair	no
<=30	high	no	excellent	no
31...40	high	no	fair	yes
>40	medium	no	fair	yes
>40	low	yes	fair	yes
>40	low	yes	excellent	no
31...40	low	yes	excellent	yes
<=30	medium	no	fair	no
<=30	low	yes	fair	yes
>40	medium	yes	fair	yes
<=30	medium	yes	excellent	yes
31...40	medium	no	excellent	yes
31...40	high	yes	fair	yes
>40	medium	no	excellent	no

# Decision Tree Induction: ID3

- Root is identified first, then the parents of next label, and so on of the decision tree are identified using entropy theory based optimization function
- Resulting tree:



age	income	student	credit_rating	buys_computer
<=30	high	no	fair	no
<=30	high	no	excellent	no
31...40	high	no	fair	yes
>40	medium	no	fair	yes
>40	low	yes	fair	yes
>40	low	yes	excellent	no
31...40	low	yes	excellent	yes
<=30	medium	no	fair	no
<=30	low	yes	fair	yes
>40	medium	yes	fair	yes
<=30	medium	yes	excellent	yes
31...40	medium	no	excellent	yes
31...40	high	yes	fair	yes
>40	medium	no	excellent	no

# Algorithm of ID3

- Basic algorithm (a greedy algorithm)
  - Tree is constructed in a **top-down recursive divide-and-conquer manner**
  - At start, all the training examples are at the root
  - Attributes are categorical (if continuous-valued, they are discretized in advance)
  - Examples are partitioned recursively based on selected attributes
  - Test attributes are selected on the basis of a heuristic or statistical measure (e.g., **information gain**)
- Conditions for stopping partitioning
  - All samples for a given node belong to the same class
  - There are no remaining attributes for further partitioning – **majority voting** is employed for classifying the leaf
  - There are no samples left

# Brief Review of Entropy

## ■ Entropy (Information Theory)

- A measure of uncertainty associated with a random variable

- Calculation: For a discrete random variable  $Y$  taking  $m$  distinct values  $\{y_1, \dots, y_m\}$ ,


- $H(Y) = -\sum_{i=1}^m p_i \log(p_i)$ , where  $p_i = P(Y = y_i)$

- Interpretation:

- Higher entropy => higher uncertainty
  - Lower entropy => lower uncertainty

## ■ Conditional Entropy


- $H(Y|X) = \sum_x p(x)H(Y|X = x)$



$Y/(X=x_1)$	$Y/(X=x_2)$	.....	$Y/(X=x)$	$Y/(X=x_n)$

# Entropy and uncertainty

- Suppose three patients have completed a medical test which, yields one of two possible results: the disease is either present or absent.
- Let Patient A has 95% chance that he has the disease. For Patient B and C it is 30% and 50%, respectively.
- So, if A, B, and C are in waiting room of a doctor's office, then the uncertainty of the waiting room is:





# Entropy and uncertainty

- All other things being equal, which of the three patients is confronted with the greatest degree of uncertainty?
- I think the answer is clear: patient C.
- Compare this with patient A. Patient A is experiencing little uncertainty with regard to his medical prospects.
- Intuitively speaking, uncertainty of patient B falls in between that of A and C.
- **Measuring uncertainty:** Entropy is a measure of uncertainty.
- Entropy allows us to make precise statements and perform computations with regard to one of life's most pressing issues: not knowing how things will turn out.
- By the term entropy, I will refer to **Shannon entropy**, which is used most frequently in natural language processing and machine learning.

# Entropy and uncertainty

- The **Shannon entropy** formula for an event  $X$  with  $n$  possible outcomes and probabilities  $p_1, \dots, p_n$ :

$$H(X) = H(p_1, \dots, p_n) = - \sum_{i=1}^n p_i \log_2 p_i$$

- For previous patients example,

$$H(A) = -0.95 \log(0.95) - 0.05 \log(0.05) = 0.08$$

$$H(B) = -0.7 \log(0.7) - 0.3 \log(0.3) = 0.27$$

$$H(C) = -0.5 \log(0.5) - 0.5 \log(0.5) = 0.3 \text{ (maximum)}$$

# Entropy and uncertainty

- The **Shannon entropy** formula for an event  $X$  with  $n$  possible outcomes and probabilities  $p_1, \dots, p_n$ :

$$H(X) = H(p_1, \dots, p_n) = - \sum_{i=1}^n p_i \log_2 p_i$$


- **Basic properties of uncertainty:**

**Prop. 1: Uniform distributions have maximum uncertainty**

If your goal is to minimize uncertainty, stay away from **uniform probability distributions**.

# Prop -1 of uncertainty

- A probability distribution is a function that assigns a probability to every possible outcome such that the probabilities add up to 1.
- A distribution is uniform when all of the outcomes have the same probability.
- For example, fair coins (50% heads, 50% tails) and fair dice (1/6 probability for each of the six faces) follow uniform distributions.




# Prop -1 of uncertainty

- A good measure of uncertainty achieves its highest values for uniform distributions.
- Entropy satisfies the criterion. Given  $n$  possible outcomes, maximum entropy is maximized by equiprobable outcomes:

$$p_1 = \dots = p_n = \frac{1}{n}$$

- Here is the plot of the Entropy function as applied to Bernoulli trials (events with two possible outcomes and probabilities  $p$  and  $1-p$ ):
- In the case of Bernoulli trials, entropy reaches its maximum value for  $p=0.5$



# Prop -2 of uncertainty

- **Prop-2: Uncertainty is additive for independent events**
- Let  $A$  and  $B$  be independent events. In other words, knowing the outcome of event  $A$  does not tell us anything about the outcome of event  $B$ .
- The uncertainty associated with both events should be the sum of the individual uncertainties:


$$H(X, Y) = H(X) + H(Y)$$

- Let's use the example of flipping two coins to make this more concrete. We can either flip both coins simultaneously or first flip one coin and then flip the other one.
- In other words, we can either report the outcome of the two coin flips at once or separately. The uncertainty is the same in either case.

# Prop -2 of uncertainty

- Ex: Let the first coin lands heads ( $H$ ) up with an 80% probability and tails ( $T$ ) up with a probability of 20%.
- The probabilities for the other coin are 60% and 40%.
- If we flip both coins simultaneously, there are four possible outcomes:  $HH$ ,  $HT$ ,  $TH$  and  $TT$ . The corresponding probabilities are given by  $[0.48, 0.32, 0.12, 0.08]$ .

- The joint entropy (green) for the two independent events is equal to the sum of the individual events (red and blue).



# Prop -2 of uncertainty

- Plugging the numbers into the entropy formula, we see that: Just as promised,


$$H(0.48, 0.32, 0.12, 0.08) = H(0.8, 0.2) + H(0.6, 0.4)$$



# Prop -3 of uncertainty

## Prop-3: Adding an outcome with zero probability has no effect

- Suppose (a) you win whenever outcome #1 occurs and (b) you can choose between two probability distributions, *A* and *B*.
- Distribution *A* has two outcomes: say, 80% and 20%. Distribution *B* has three outcomes with probabilities 80%, 20% and 0%.



# Prop -3 of uncertainty

- Given the options  $A$  and  $B$ , which one would you choose?
- Ans: The inclusion of the third outcome neither increases nor decreases the uncertainty associated with the game.  $A$  or  $B$ , who cares. It doesn't matter.
- The entropy formula agrees with this assessment:
$$H(p_1, \dots, p_n) = H(p_1, \dots, p_n, 0)$$
- In words, adding an outcome with zero probability has no effect on the measurement of uncertainty.

# Prop -4 of uncertainty

## **Prop-4: The measure of uncertainty is continuous in all its arguments**

- The last of the basic properties is continuity.
- Famously, the intuitive explanation of a continuous function is that arbitrarily small changes in the output (uncertainty, in our case) should be achievable through sufficiently small changes in the input (probabilities).
- Logarithm functions are continuous at every point for which they are defined. So are sums and products of a finite number of functions that are continuous on a subset.
- It follows that the entropy function is continuous in its probability arguments.

# Entropy

- **The Uniqueness Theorem:** [Khinchin \(1957\)](#) showed that the only family of functions satisfying the four basic properties described above is of the following form:

$$H(p_1, \dots, p_n) = -\lambda \sum_{i=1}^n p_i \log p_i,$$

- where  $\lambda$  is a positive constant. Khinchin referred to this as the **Uniqueness Theorem**.
- Setting  $\lambda = 1$  and using the binary logarithm gives us the Shannon entropy.
- To reiterate, entropy is used because it has desirable properties and is the natural choice among the family functions that satisfy all items on the basic properties.

# Entropy

## Other properties of Entropy:

- Entropy has many other properties used in Khinchin's Uniqueness Theorem. Some of them:

### Prop-5: Uniform distributions with more outcomes have more uncertainty


- Suppose you have the choice between a fair coin and a fair die:

Fair coin or fair die ?

Let, A for coin and B for die.

$$H(A) = -0.5\ln(0.5) - 0.5\ln(0.5) = 0.69$$

$$H(B) = [-(1/6)\ln(1/6)] \times 6 = 1.8$$



# Entropy

- And let's say you win if the coin lands heads up or the die lands on face 1.
- Which of the two options would you choose?
  - (i) if you are a profit maximizer and
  - (ii) if you prefer with more variety and uncertainty.
- As the number of equiprobable outcomes increases, so should our measure of uncertainty.
- And this is exactly what Entropy does:  $H(1/6, 1/6, 1/6, 1/6, 1/6, 1/6) > H(0.5, 0.5)$ .
- And, in general, if we let  $L(k)$  be the entropy of a uniform distribution with  $k$  possible outcomes,

$$L(m) > L(n) \quad \text{for } m > n$$


# Entropy

## Prop-6: Events have non-negative uncertainty

- Do you know what negative uncertainty is? Neither do I.
- A user-friendly measure of uncertainty should always return a non-negative quantity, no matter what the input is.
- This is yet another criterion that is satisfied by entropy. Let's take another look at the formula:  $H(X) = - \sum_{i=1}^n p_i \log_2 p_i$
- Probabilities are, by definition, in the range between 0 and 1 and, therefore, non-negative.
- The logarithm of a probability is non-positive. Multiplying the logarithm of a probability with a probability doesn't change the sign. The sum of non-positive products is non-positive. And finally, the negative of a non-positive value is non-negative.
- Entropy is, thus, non-negative for every possible input.

# Entropy

- **Prop-7: Events with a certain outcome have zero uncertainty**
- Suppose you are in possession of a magical coin. No matter how you flip always lands head up.
  - Suppose that outcome  $i$  certain to occur. It follows that  $p_i$ , the probability of outcome  $i$ , is equal to




1.  $H(X)$ , thus, simplifies to:

$$H(0, \dots, 1, \dots, 0) = -\log_2 1 = 0$$



# Entropy

- **Prop-8: Flipping the arguments has no effect**
- This is another obviously desirable property. Consider two cases. In the first case, the probability of heads and tails are 80% and 20%. In the second case, the probabilities are reversed: heads 20%, tails 80%.



# Entropy

- Both coin flips are equally uncertain and have the same entropy:  $H(0.8, 0.2) = H(0.2, 0.8)$ .
- In more general terms, for the case of two outcomes, we have:

$$H(p_1, p_2) = H(p_2, p_1)$$

- This fact applies to any number of outcomes. We can position the arguments (i.e., the probabilities of a distribution) in any order we like. The result of the entropy function is always the same.

# Attribute Selection Measure: Information Gain (ID3/C4.5)

- Select the attribute with the highest information gain
- Let  $p_i$  be the probability that an arbitrary tuple in  $D$  belongs to class  $C_i$ , estimated by  $|C_{i,D}|/|D|$
- Entropy (a measure of unpredictability or impurity) of the class or target variable in the dataset = Expected information needed to classify a tuple in  $D$ :

$$Info(D) = -\sum_{i=1}^m p_i \log_2(p_i)$$

- Entropy of the target variable after split the dataset  $D$  based on a feature  $A$  = Weighted average of the entropies of the subsets formed by the split of  $D$  using  $A$  = Information needed (after using  $A$  to split  $D$  into  $v$  partitions) to classify a tuple in  $D$ :

$$Info_A(D) = \sum_{j=1}^v \frac{|D_j|}{|D|} \times Info(D_j)$$

# Attribute Selection Measure: Information Gain (ID3/C4.5)

- The information gain is the difference between the entropy before the split and the entropy after the split.

$$Gain(A) = Info(D) - Info_A(D)$$

- Information gain is the reduction in entropy after a dataset is split on a feature.
- A higher information gain indicates that the feature provides a better separation of the target variable, making it a preferred choice for splitting the dataset in decision tree algorithms.

# Attribute Selection: Information Gain

■ Class P: buys\_computer = “yes”

■ Class N: buys\_computer = “no”

$$Info_{age}(D) = \frac{5}{14} I(2,3) + \frac{4}{14} I(4,0) + \frac{5}{14} I(3,2) = 0.694$$

$$Info(D) = I(9,5) = -\frac{9}{14} \log_2\left(\frac{9}{14}\right) - \frac{5}{14} \log_2\left(\frac{5}{14}\right) = 0.940$$

age	p <sub>i</sub>	n <sub>i</sub>	I(p <sub>i</sub> , n <sub>i</sub> )
<=30	2	3	0.971
31...40	4	0	0
>40	3	2	0.971

$\frac{5}{14} I(2,3)$  means “age <=30” has 5 out of 14 samples, with 2 yes’ es and 3 no’ s. Hence

$$Gain(age) = Info(D) - Info_{age}(D) = 0.246$$

Similarly,

$$Gain(income) = 0.029$$

$$Gain(student) = 0.151$$

$$Gain(credit\_rating) = 0.048$$

age	income	student	credit_rating	buys_computer
<=30	high	no	fair	no
<=30	high	no	excellent	no
31...40	high	no	fair	yes
>40	medium	no	fair	yes
>40	low	yes	fair	yes
>40	low	yes	excellent	no
31...40	low	yes	excellent	yes
<=30	medium	no	fair	no
<=30	low	yes	fair	yes
>40	medium	yes	fair	yes
<=30	medium	yes	excellent	yes
31...40	medium	no	excellent	yes
31...40	high	yes	fair	yes
>40	medium	no	excellent	no

# Computing Information-Gain for Continuous-Valued Attributes

- Let attribute A be a continuous-valued attribute
- Must determine the *best split point* for A
  - Sort the value A in increasing order
  - Typically, the midpoint between each pair of adjacent values is considered as a possible *split point*
    - $(a_i + a_{i+1})/2$  is the midpoint between the values of  $a_i$  and  $a_{i+1}$
  - The point with the *minimum expected information requirement* for A is selected as the split-point for A
- Split:
  - D1 is the set of tuples in D satisfying  $A \leq \text{split-point}$ , and D2 is the set of tuples in D satisfying  $A > \text{split-point}$

# Gain Ratio for Attribute Selection (C4.5)

- Information gain measure is biased towards attributes with a large number of values
- C4.5 (a successor of ID3) uses gain ratio to overcome the problem (normalization to information gain)

$$SplitInfo_A(D) = - \sum_{j=1}^v \frac{|D_j|}{|D|} \times \log_2 \left( \frac{|D_j|}{|D|} \right)$$

$$- \text{GainRatio}(A) = \text{Gain}(A) / \text{SplitInfo}(A)$$

- Ex.  $SplitInfo_{income}(D) = -\frac{4}{14} \times \log_2 \left( \frac{4}{14} \right) - \frac{6}{14} \times \log_2 \left( \frac{6}{14} \right) - \frac{4}{14} \times \log_2 \left( \frac{4}{14} \right) = 1.557$

$$- \text{gain\_ratio}(\text{income}) = 0.029 / 1.557 = 0.019$$

$$\text{Similarly, Gain\_ratio}(\text{student}) = 0.151 / 0.69 = 0.22$$

- The attribute with the maximum gain ratio is selected as the splitting attribute

# Gini Index

- The Gini coefficient measures the [inequality](#) among the values of a [frequency distribution](#), such as levels of [income](#). In machine learning, it is utilized as an impurity measure in decision tree algorithms for classification tasks.
- The Gini index is the most commonly used measure of inequality. It was developed by Italian statistician Corrado Gini and is named after him.
- A Gini coefficient of 0 reflects perfect equality, where all income or wealth values are the same, while a Gini coefficient of 1 (or 100%) reflects maximal inequality among values, a situation where a single individual has all the income while all others have none.



# Gini Index (CART, IBM IntelligentMiner)

- If a data set  $D$  contains examples from  $n$  classes, gini index,  $gini(D)$  is defined as

$$gini(D) = 1 - \sum_{j=1}^n p_j^2$$

where  $p_j$  is the relative frequency of class  $j$  in  $D$

- If a data set  $D$  is split on  $A$  into two subsets  $D_1$  and  $D_2$ , the  $gini$  index  $gini_A(D)$  is defined as

$$gini_A(D) = \frac{|D_1|}{|D|} gini(D_1) + \frac{|D_2|}{|D|} gini(D_2)$$

- Reduction in Impurity:

$$\Delta gini(A) = gini(D) - gini_A(D)$$

- The attribute provides the largest reduction in impurity is chosen to split the node.

# Computation of Gini Index

- Ex. D has 9 tuples in buys\_computer = “yes” and 5 in “no”

$$gini(D) = 1 - \left(\frac{9}{14}\right)^2 - \left(\frac{5}{14}\right)^2 = 0.459$$

- Suppose the attribute income partitions D into 10 in  $D_1$ : {low, medium} and 4 in  $D_2$   $gini_{income \in \{low, medium\}}(D) = \left(\frac{10}{14}\right)Gini(D_1) + \left(\frac{4}{14}\right)Gini(D_2)$

$$= \frac{10}{14} \left(1 - \left(\frac{7}{10}\right)^2 - \left(\frac{3}{10}\right)^2\right) + \frac{4}{14} \left(1 - \left(\frac{2}{4}\right)^2 - \left(\frac{2}{4}\right)^2\right)$$

$$= 0.443$$

$$= Gini_{income \in \{high\}}(D).$$

$Gini_{\{low, high\}}$  is 0.458;  $Gini_{\{medium, high\}}$  is 0.450. Thus, split on the {low, medium} (and {high}) since it has the lowest Gini index

age	income	student	credit_rating	buys_computer
<=30	high	no	fair	no
<=30	high	no	excellent	no
31...40	high	no	fair	yes
>40	medium	no	fair	yes
>40	low	yes	fair	yes
>40	low	yes	excellent	no
31...40	low	yes	excellent	yes
<=30	medium	no	fair	no
<=30	low	yes	fair	yes
>40	medium	yes	fair	yes
<=30	medium	yes	excellent	yes
31...40	medium	no	excellent	yes
31...40	high	yes	fair	yes
>40	medium	no	excellent	no

# Comparing Attribute Selection Measures

- The three measures, in general, return good results but
  - **Information gain:**
    - biased towards multivalued attributes
  - **Gain ratio:**
    - tends to prefer unbalanced splits in which one partition is much smaller than the others
  - **Gini index:**
    - biased to multivalued attributes
    - has difficulty when # of classes is large

# Other Attribute Selection Measures

- CHAID: a popular decision tree algorithm, measure based on  $\chi^2$  test for independence
- C-SEP: performs better than info. gain and gini index in certain cases
- G-statistic: has a close approximation to  $\chi^2$  distribution
- MDL (Minimal Description Length) principle (i.e., the simplest solution is preferred):
  - The best tree as the one that requires the fewest # of bits to both (1) encode the tree, and (2) encode the exceptions to the tree
- Multivariate splits (partition based on multiple variable combinations)
  - CART: finds multivariate splits based on a linear comb. of attrs.
- Which attribute selection measure is the best?
  - Most give good results, none is significantly superior than others

# Overfitting and Tree Pruning

- Overfitting: An induced tree may overfit the training data
  - Too many branches, some may reflect anomalies due to noise or outliers
  - Poor accuracy for unseen samples
- Two approaches to avoid overfitting
  - Prepruning: *Halt tree construction early*-do not split a node if this would result in the goodness measure falling below a threshold
    - Difficult to choose an appropriate threshold
  - Postpruning: *Remove branches* from a “fully grown” tree—get a sequence of progressively pruned trees
    - Use a set of data different from the training data to decide which is the “best pruned tree”